



VISHNU A S

GENERATIVE AI ENGINEER

✉ Vishnuashok9811@gmail.com ☎ +91 8122455847 📍 Bengaluru, India  [linkedin.com/in/vishnu-as-7734a11ab/](https://www.linkedin.com/in/vishnu-as-7734a11ab/)

Professional Summary

- Generative AI Engineer with 4.5 years of experience designing, developing, and deploying end-to-end LLM-powered applications, including chatbots, document Q&A systems, and enterprise knowledge assistants built on Retrieval Augmented Generation (RAG) architecture.
- Skilled in prompt engineering and prompt optimization, developing reusable and version-controlled prompt templates that improved response accuracy and reduced hallucination rates in production AI systems.
- Hands-on experience designing and implementing RAG pipelines, integrating Large Language Models with vector databases such as Pinecone and FAISS to deliver high-precision semantic search.
- Proficient in integrating LLMs through REST APIs, including Claude models via AWS Bedrock and Azure Open AI.
- Experienced in building AI workflow automation solutions that streamline document processing, summarization, and enterprise search across PDF, CSV, and relational database sources.
- Strong background in vector database implementation and embedding generation using Sentence Transformers and modern embedding models to power efficient similarity search.
- Deployed scalable, serverless Generative AI applications on AWS using Lambda, API Gateway, S3, and Step Functions, ensuring high availability and cost-efficient operations.
- Implemented logging, monitoring, and performance optimization practices for AI pipelines, improving system reliability and reducing average LLM response latency.
- Experience handling enterprise knowledge bases, unifying structured and unstructured data sources into consolidated Gen AI-powered search and Q&A systems.
- Solved real-world business problems using Generative AI, including automating document review, reducing manual search effort, and improving decision-making through AI-driven insights.
- Familiar with prompt versioning, model versioning, and responsible AI practices, ensuring safe, controlled, and auditable AI system behavior in enterprise production environments.

GENERATIVE AI ENGINEER

Hcl Tech, Bengaluru, India

Feb 2025 – Present

Banking Fraud Alert System (AWS)

Client: Confidential Client (Banking & Financial Services) | Domain: BFSI

- Designed and developed a Generative AI-powered fraud alert system for banking transactions, using LLMs to analyze suspicious activity and generate clear, investigator-ready alert summaries.
- Integrated an internally developed claude with a Retrieval Augmented Generation (RAG) pipeline to contextualize flagged transactions against historical fraud patterns and compliance policies.

EDUCATION

BACHELOR OF ENGINEERING

Anna University, Tamil Nadu, India | 2020

KEY SKILLS

- Python
- Generative AI
- Agentic AI
- Large Language Models (LLMs)
- Retrieval-Augmented Generation (RAG)
- LangChain
- LangGraph
- Prompt Engineering
- AI Agents
- Multi-Agent Systems
- FastAPI
- RESTful APIs
- AWS Bedrock
- AWS Lambda
- Pinecone
- FAISS
- OpenSearch
- Vector Databases
- Semantic Search
- Embedding Models
- OpenAI GPT-4
- OCR (AWS Textract & Tesseract)
- LangSmith
- Chunking Strategies
- Conversational AI
- Tool Calling
- Pydantic
- Metadata Filtering
- RAGAS
- Jira

- ♦ Built and maintained a FAISS-based vector database to store embeddings of historical fraud cases, enabling fast similarity search and pattern matching for new alerts.
- ♦ Generated embeddings for transaction records and case documentation using SentenceTransformers, improving retrieval precision for fraud-pattern identification.
- ♦ Developed Langgraph-based pipelines to orchestrate prompt chains, contextual data retrieval, and LLM-driven alert generation.
- ♦ Exposed fraud detection and alerting capabilities through REST APIs for integration with core banking fraud monitoring systems.
- ♦ Deployed and hosted pipeline components on AWS (Lambda, S3, API Gateway), ensuring high availability and low-latency alert generation at scale.
- ♦ Implemented validation and review checks on AI-generated alerts to minimize false positives and maintain regulatory compliance.

Senior Software Engineer

LTIMindtree Chennai, India

Mar 2022 – Jan 2025

Healthcare Claim Adjudication Platform (Azure)

Client: Confidential Client (Healthcare) | Domain: Healthcare

- ♦ Designed and developed a Generative AI-powered claim adjudication assistant using Azure OpenAI Service to automate the review and decisioning of healthcare insurance claims.
- ♦ Built RAG pipelines using Langgraph, integrating Azure OpenAI models with Azure AI Search to retrieve relevant policy documents and claim history.
- ♦ Implemented embedding models to convert claim forms, medical records, and policy documents into vector representations for accurate, efficient retrieval.
- ♦ Developed serverless AI workflows using Azure Functions, exposing endpoints through Azure API Management for consumption by internal claims processing systems.
- ♦ Orchestrated multi-step claim adjudication workflows using Azure Logic Apps, automating document ingestion, embedding generation, and adjudication recommendations end-to-end.
- ♦ Managed document storage and retrieval using Azure Blob Storage, integrating PDF, CSV, and relational database sources into a unified adjudication pipeline.
- ♦ Applied prompt engineering techniques to optimize claim adjudication recommendations, improving accuracy while reducing token usage and cost.
- ♦ Implemented monitoring and logging using Azure Monitor, tracking pipeline performance and identifying opportunities for further optimization.

ADDITIONAL INFO

Date of birth : 22.11.1998
Nationality : Indian
Gender : Male
Marital Status : Single

LANGUAGES

- ♦ **English**
Advanced knowledge
- ♦ **Tamil**
Working Knowledge
- ♦ **Malayalam**
Native